

Com S // CPR E // MATH 5250

Numerical Analysis of High-Performance Computing

Instructor: Songting Luo

Lecture 3: Floating Point Numbers and More Python

Outline

1. Announcements
2. Floating Point Numbers
3. More Python

Announcements

Installing relevant software on a personal computer

- All software used in this class is open-source (with the exception of software on the HPC-Class cluster)
- Therefore, it is possible for you to setup your own personal machine
- On a **Linux** machine, installation of **gcc**, **git**, **python**, **IPython**, **matplotlib**, **pylab** is either automatic or easy
- On a **Windows** machine: Install VirtualBox:
<https://www.virtualbox.org/>

Or Windows Subsystem for Linux (WSL):

<https://docs.microsoft.com/en-us/windows/wsl/>

This will let you run Linux inside of Windows.

- On a **Mac OS X** machine:
 - Install Developer's Tools (Xcode, etc...) from <http://www.apple.com>
 - Install **gcc** and **git** via homebrew (<http://brew.sh/>) or Fink (<http://www.finkproject.org/>)
 - Install **python**, **IPython**, **matplotlib**, **pylab**

The power of using Git

- No matter what computer you are on (Carver 449 Lab, your own laptop, another computer lab, etc...), you can always work on your Math 5250 Labs by **cloning** your repository to your current machine:

```
git clone {repository on GitHub}
```

- git push from local machine to GitHub

```
$ git add -A .  
$ git commit -m "whatever message to keep track of your files"  
$ git push
```

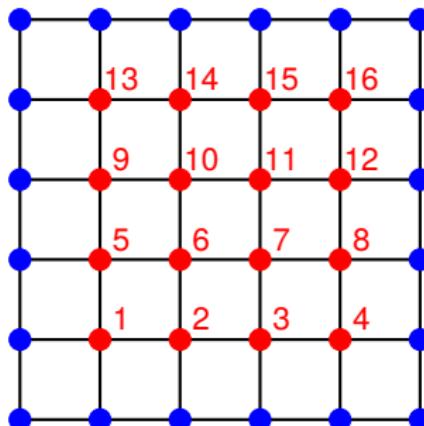
- git pull from GitHub to local machine

```
$ git pull
```

Floating Point Numbers

Recall: steady-state heat conduction

- Discretize an $N \times N$ grid with N^2 unknowns



- Assume temperature is fixed (and known) at each **boundary point**
- neighboring value = node that shares an edge with current node
- Every **interior point** has **exactly** 4 neighbors
- For each **interior point**, the steady state value is approximately the average of the 4 neighboring values

Steady-state heat conduction

System size: If $N = 100 \implies 100 \times 100$ grid \implies linear system with 10,000 equations, $A\vec{U} = \vec{b}$, where A is of size $10,000 \times 10,000$.

Question: How much disk space is required to store a $10,000 \times 10,000$ matrix of real numbers?

Steady-state heat conduction

System size: If $N = 100 \implies 100 \times 100$ grid \implies linear system with 10,000 equations, $A\vec{U} = \vec{b}$, where A is of size $10,000 \times 10,000$.

Question: How much disk space is required to store a $10,000 \times 10,000$ matrix of real numbers?

Answer: That depends on how many bytes are used for each real number.
1 byte = 8 bits, bit = "binary digit"

Steady-state heat conduction

System size: If $N = 100 \implies 100 \times 100$ grid \implies linear system with 10,000 equations, $A\vec{U} = \vec{b}$, where A is of size $10,000 \times 10,000$.

Question: How much disk space is required to store a $10,000 \times 10,000$ matrix of real numbers?

Answer: That depends on how many bytes are used for each real number.
1 byte = 8 bits, bit = “binary digit”

Assuming 8 bytes (64 bits) per value:

$$\begin{aligned}10,000 \times 10,000 \text{ matrix} &\implies 10^8 \text{ entries matrix} \\ \implies 8 \times 10^8 \text{ bytes} &= 8 \times 10^5 \text{ kB} = 8 \times 10^2 \text{ MB} = \mathbf{800 \text{ MB}}\end{aligned}$$

Steady-state heat conduction

System size: If $N = 100 \implies 100 \times 100$ grid \implies linear system with 10,000 equations, $A\vec{U} = \vec{b}$, where A is of size $10,000 \times 10,000$.

Question: How much disk space is required to store a $10,000 \times 10,000$ matrix of real numbers?

Answer: That depends on how many bytes are used for each real number.
1 byte = 8 bits, bit = “binary digit”

Assuming 8 bytes (64 bits) per value:

$$\begin{aligned}10,000 \times 10,000 \text{ matrix} &\implies 10^8 \text{ entries matrix} \\ \implies 8 \times 10^8 \text{ bytes} &= 8 \times 10^5 \text{ kB} = 8 \times 10^2 \text{ MB} = \textcolor{red}{800 \text{ MB}}\end{aligned}$$

Note: less than 50,000 values are non-zero \implies **99.95% are 0.**

Therefore, if we can get away with only storing non-zeros:

$$800 \text{ MB} \times 0.0005 = 0.4 \text{ MB} = \textcolor{red}{400 \text{ kB}}$$

Measuring size and speed

Kilo = thousand (10^3)

Mega = million (10^6)

Giga = billion (10^9)

Tera = trillion (10^{12})

Peta = 10^{15}

Exa = 10^{18}

Binary representation

Memory is divided into **bytes**, consisting of 8 bits each.

1 byte can hold $2^8 = 256$ distinct numbers:

00000000 = 0

00000001 = 1

00000010 = 2

...

11111111 = 255

Binary representation

Memory is divided into **bytes**, consisting of 8 bits each.

1 byte can hold $2^8 = 256$ distinct numbers:

$$00000000 = 0$$

$$00000001 = 1$$

$$00000010 = 2$$

...

$$11111111 = 255$$

8-bit representation: $(a_7)(a_6)(a_5)(a_4)(a_3)(a_2)(a_1)(a_0)$

Converting a binary representation to an integer: $x = \sum_{n=0}^7 a_n 2^n$

$$11111111 : x = \sum_{n=0}^7 2^n = 1 + 2 + 4 + 8 + 16 + 32 + 64 + 128 = 255$$

A **byte** might represent integers, characters, colors, ...

Modern computers: 4 bytes (32 bits) or 8 bytes (64 bits) used to represent integers and real numbers.

Integers

To store integers, need one bit for the sign (+ or -). In one byte, this leaves 7 bits.

Two's complement representation:

10000000 = -128

10000001 = -127

10000010 = -126

...

11111110 = -2

11111111 = -1

00000000 = 0

00000001 = 1

00000010 = 2

...

01111111 = 127

$$x = -a_7 2^7 + \sum_{n=0}^6 a_n 2^n$$

Integers

To store integers, need one bit for the sign (+ or -). In one byte, this leaves 7 bits.

Two's complement representation:

$$10000000 = -128$$

$$10000001 = -127$$

$$10000010 = -126$$

...

$$11111110 = -2$$

$$11111111 = -1$$

$$00000000 = 0$$

$$00000001 = 1$$

$$00000010 = 2$$

...

$$01111111 = 127$$

$$x = -a_7 2^7 + \sum_{n=0}^6 a_n 2^n$$

Advantage: binary addition works directly (e.g., $10 + (-6) = 4$)

$$\begin{array}{r} & 0 & 0 & 0 & 0 & 1 & 0 & 1 & 0 \\ + & 1 & 1 & 1 & 1 & 1 & 0 & 1 & 0 \\ \hline [1] & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 \end{array}$$

$$\implies 00000100 = 2^2 \text{ (base 10)} = 4 \text{ (base 10)}$$

Integers

Integers are typically stored in 8 bytes (64 bits). Values roughly between

$$-2^{63} \approx -10^{19} \quad \text{and} \quad 2^{63} \approx 10^{19}$$

can be stored.

In Python, larger integers will automatically be stored using more bytes.

```
$ python
```

```
>> 2**30
1073741824
```

```
>>> 2**62
4611686018427387904
```

```
>>> 2**63
9223372036854775808L
```

```
>>> 2**100
1267650600228229401496703205376L
```

Fixed point numbers

Idea: Use 64 bits for a real number but always assume N bits in integer part and M bits in fractional part.

Example: 5 digits for integer part and 6 digits in fractional part:

| | |
|--------------|--------------|
| 00003.141592 | (pi) |
| 00000.000314 | (pi / 10000) |
| 31415.926535 | (pi * 10000) |

Fixed point numbers

Idea: Use 64 bits for a real number but always assume N bits in integer part and M bits in fractional part.

Example: 5 digits for integer part and 6 digits in fractional part:

| | |
|--------------|--------------|
| 00003.141592 | (pi) |
| 00000.000314 | (pi / 10000) |
| 31415.926535 | (pi * 10000) |

Disadvantages:

- Precision depends on size of number
- Often many wasted bits (leading 0's)
- Limited range: often scientific problems involve very large or small numbers

Floating point numbers

Base 10 scientific notation:

$$0.31278e-8 = 0.31278 \times 10^{-8} = 0.000000031278$$

Mantissa: 0.31278, Exponent: - 8

Floating point numbers

Base 10 scientific notation:

$$0.31278e-8 = 0.31278 \times 10^{-8} = 0.000000031278$$

Mantissa: 0.31278, Exponent: - 8

Binary floating point numbers:

Example:

Mantissa: 0.101101, Exponent: - 11011

means

$$\begin{aligned}0.101101 &= 1 \times (2^{-1}) + 0 \times (2^{-2}) + 1 \times (2^{-3}) + 1 \times (2^{-4}) + 0 \times (2^{-5}) + 1 \times (2^{-6}) \\&= 0.703125 \text{ (base 10)}\end{aligned}$$

$$\begin{aligned}-11011 &= - \left\{ 1 \times (2^4) + 1 \times (2^3) + 0 \times (2^2) + 1 \times (2^1) + 1 \times (2^0) \right\} \\&= -27 \text{ (base 10)}\end{aligned}$$

So the number is

$$0.703125 \times 2^{-27} \approx 5.238689482212067 \times 10^{-9}$$

Floating point numbers

- Python `float` is 8 bytes with IEEE standard representation
- 53 bits for mantissa and 11 bits for exponent (64 bits = 8 bytes)
- We can store 52 binary bits of `precision`:

$$\varepsilon_{\text{machine}} = 2^{-52} \approx 2.22 \times 10^{-16}$$

- Therefore, roughly **15 digits of precision**
- Machine epsilon ($\varepsilon_{\text{machine}}$) is the relative distance between any two consecutive numbers on the floating point number system
- Therefore, simply representing a real number in a floating point number system produces an error: **round-off error**

Floating point numbers

- Since $2^{-52} \approx 2.22 \times 10^{-16}$, this corresponds to roughly 15 digits of precision
- For example:

```
>>> from numpy import pi
```

```
>>> pi
```

```
3.141592653589793
```

```
>>> 1000*pi
```

```
3141.592653589793
```

```
>>> pi/1000
```

```
0.0031415926535897933
```

- Note: storage and arithmetic is done in base 2; converted to base 10 only when printed!

More Python

Python overview

Python is an object oriented general-purpose language

Advantages:

- Can be used interactively from a Python shell (similar to Matlab)
- Can also write scripts to execute from Unix shell
- Little overhead to start programming
- Powerful modern language
- Many modules are available for specialized work
- Good graphics and visualization modules (e.g., Matplotlib)
- Easy to combine with other languages (e.g., C and Fortran)
- Open source and runs on all platforms

Python overview

Disadvantage: Can be slow to do certain things, such as looping over arrays

This is because Python code is **interpreted** rather than **compiled**

Additional features to improve performance:

- Need to use suitable modules (e.g., NumPy, Pandas) for speed
- Can easily create custom modules from compiled code written in Fortran, C, etc...
- Can also use extensions such as **Cython** that makes it easier to mix Python with C code that will be compiled
- Python is often used for high-level scripts that e.g., download data from the web, run a set of experiments, collate and plot results

Python is an object-oriented language

- Nearly everything in Python is an object of some class
- The class description tells what data the object holds (attributes) and what operations (methods or functions) are defined to interact with the object
- Every “**variable**” is really just a pointer to some object. You can reset it to point to some other object at will
- So variables don’t have “**type**” (e.g., integer, float, string) (but the objects they currently point to do)

Python is an object-oriented language

- Nearly everything in Python is an object of some class
- The class description tells what data the object holds (attributes) and what operations (methods or functions) are defined to interact with the object
- Every “**variable**” is really just a pointer to some object. You can reset it to point to some other object at will
- So variables don’t have “**type**” (e.g., integer, float, string) (but the objects they currently point to do)

```
>>> team = 'Cyclones'  
>>> print(team)  
Cyclones  
>>> rival = 'Hawkeyes'  
>>> t = team  
>>> print(t)  
Cyclones
```

| variable | value |
|----------|------------|
| team | 'Cyclones' |
| rival | 'Hawkeyes' |
| t | |

A diagram illustrating variable pointers. It consists of a table with two columns: 'variable' and 'value'. The 'variable' column contains three entries: 'team', 'rival', and 't'. The 'value' column contains three entries: "'Cyclones'", "'Hawkeyes'", and an empty string. Red arrows point from each variable name to its corresponding value. A red 'X' is drawn over the entry for 'rival' in both the variable and value columns. Additionally, a red 'X' is drawn over the entry for 't' in the variable column.

Python is an object-oriented language

- Nearly everything in Python is an object of some class
- The class description tells what data the object holds (attributes) and what operations (methods or functions) are defined to interact with the object
- Every “**variable**” is really just a pointer to some object. You can reset it to point to some other object at will
- So variables don’t have “**type**” (e.g., integer, float, string) (but the objects they currently point to do)

```
>>> team = 'Cyclones'  
>>> print(team)  
Cyclones  
>>> rival = 'Hawkeyes'  
>>> t = team  
>>> print(t)  
Cyclones  
>>> t = 4.0**3  
>>> print(t)  
64.0
```

| variable | value |
|----------|------------|
| team | 'Cyclones' |
| rival | 'Hawkeyes' |
| t | |

Python is an object-oriented language

```
>>> x = -72.05
>>> print(id(x), type(x))
140485917044336 <class 'float'>

>>> x = 11.83
>>> print( id(x), type(x) )
140485886550992 <class 'float'>

>>> x = [-2,0,2]
>>> print( id(x), type(x) )
140485916373568 <class 'list'>

>>> x = [7,3,-1]
>>> print( id(x), type(x) )
140485915342720 <class 'list'>
```

Python is an object-oriented language

```
>>> x = [7,3,-1]

>>> print( id(x), type(x) )
140485915342720 <class 'list'>

>>> x.append(-17)

>>> x
[7, 3, -1, -17]

>>> print( id(x), type(x) )
140485915342720 <class 'list'>
```

Note: Object of type 'list' has a method 'append' that changes the object

A list is a **mutable object** (i.e., can be changed)

Python is an object-oriented language

```
>>> x = [4,-5,7]

>>> print( id(x), x )
140485917375872 [4, -5, 7]

>>> y = x

>>> print( id(y), y )
140485917375872 [4, -5, 7]

>>> y.append(-11)

>>> y
[4, -5, 7, -11]

>>> x
[4, -5, 7, -11]
```

Note: x and y point to the same object !!!

Make a copy of a list

```
>>> x = [5,6,7]

>>> print( id(x), x )
140485915342720 [5, 6, 7]

>>> y = list(x)    # creates a new list object (contrast y = x)

>>> print( id(y), y )
140485917393920 [5, 6, 7]

>>> y.append(-3)

>>> y
[5, 6, 7, -3]

>>> x
[5, 6, 7]
```

Integers and floats are immutable

```
>>> x = 8.45
```

```
>>> print( id(x), x )  
140485886550544 8.45
```

```
>>> y = x
```

```
>>> print( id(y), y )  
140485886550544 8.45
```

```
>>> y = y + 1
```

```
>>> print( id(y), y )  
140485886550992 9.45
```

```
>>> print( id(x), x )  
140485886550544 8.45
```

Note: in order to change y, Python created a new object and left the old object intact (x still points to that object)

Lists

Note: The elements of a list can be **any objects** (need not be same type):

```
>>> L = [3, 4.5, 'abc', [1,2]]
```

```
>>> L[0]
```

```
3
```

```
>>> L[1]
```

```
4.5
```

```
>>> L[2]
```

```
'abc'
```

```
>>> L[3]
```

```
[1, 2]
```

```
>>> L[3][0]    # element 0 of L[3]
```

```
1
```

Lists

Lists have several built-in methods ([append](#), [insert](#), [sort](#), [pop](#), [reverse](#), [remove](#), ...)

```
>>> L = [3, 4.5, 'abc', [1,2]]
```

```
>>> L2 = L.pop(2)
```

```
>>> L2  
'abc'
```

```
>>> L  
[3, 4.5, [1, 2]]
```

Note: L still points to the same object, but it has changed.

Lists and tuples

```
>>> L = [3, 4.5, 'abc']  
  
>>> L[0] = 'xy'  
  
>>> L  
['xy', 4.5, 'abc']  
  
>>> T = (3, 4.5, 'abc')  
  
>>> T[0]  
3  
  
>>> T[0] = 'xy'  
Traceback (most recent call last):  
  File "<stdin>", line 1, in <module>  
TypeError: 'tuple' object does not support item assignment
```

Note: a tuple is like a list but is **immutable**

Lab assignment

- Practice operations on a Python list: `append()`, `clear()`, `copy()`, `count()`, `extend()`, `index()`, `insert()`, `pop()`, `remove()`, `reverse()`, `sort()`.
- Generate a random list L of length N compare the performance of the following operations by measuring times with increasing N : N , $2N$, $4N$, $8N$, ..., explain what you observe.
 1. Cost of removing elements from lists:
`L.pop()`
`L.pop(0)`
 2. Slicing vs explicit copy of a list:
`A = L[:]`
`B = list(L)`
 3. In-place vs out-of-place modification:
`L.reverse()`
`R = L[::-1]`
- Update your git repository by creating a folder of lecture 03, adding files and pushing to GitHub (GitLab, Bitbucket, ...)
- Submit both code and screen shots